

PaperTrail Product Specification v1.0

Overview

PaperTrail is a document intelligence platform. Users upload PDFs, notes and screenshots; the platform extracts text, builds a searchable index and answers questions with citations back to the exact page of the source document.

Core architecture

The system is composed of seven services deployed on Zerops: a static frontend (Next.js on managed Nginx), a Fastify API runtime, a Node worker that performs OCR and text extraction, a PostgreSQL database, a Valkey job queue and cache, a Meilisearch full-text index and an S3-compatible object store for originals.

The worker consumes jobs from the Valkey queue using a blocking BRPOP call. Each document is chunked into overlapping fragments that never cross page boundaries, so every search hit keeps an exact page citation.

Pricing and refunds

PaperTrail offers a free tier with 3 collections and 50 documents per month. The Pro plan costs 9 euros per month and includes unlimited collections, OCR for image documents and priority processing.

Refund policy

All paid plans are covered by a 30-day money-back guarantee from the date of purchase. Refunds are processed within 5 business days of the request and are returned to the original payment method. Annual plans are refunded pro-rata for the unused portion after the first 30 days.

Data retention

Documents are stored in object storage with AES-256 encryption at rest. Users may export or permanently delete their data at any time from the settings page.